654: Machine Learning Final Project

Havisha Khurana

Material for the Final Project

The code for the final project is spread over the following Kaggle notebooks.

1. Read image data and extract features for 50,000 images. **here**

2. Reduce features by removing zero variance, near zero variance, and medium to highly correlated features from the training data set. **here**

3. Logistic regression models on large subset (49921 rows and 1373 columns). **here**

4. Logistic regression models on small subset with all features (5004 rows and 8194 columns). **here**

5. Github repo for the project. **here**

Note: The original image data was accessed through Kaggle. See here

Research problem

The project focuses on an image classification task, precisely whether cancer was found or not in medical images of regions of the breast. Cancer is the second leading cause of deaths among American women of all races and origin (Heron 2021), and the leading cause of death in the world (Ferlay et al. 2021). Furthermore, breast cancer is among the most common forms of cancer, accounting for 10 million cases worldwide in 2020 (Ferlay et al. 2021). There are two primary types of breast cancers, Ductal carcinoma in situ (DCIS) and Invasive breast cancer (ILS or IDC), the latter making up for 70-80% of all breast cancer cases (Sharma et al. 2010). The diagnosis of the disease is time consuming and challenging process primarily because pathologist have to scan large swathes of benign regions to ultimately identify the areas of malignancy (Janowczyk and Madabhushi 2016; Cruz-Roa et al. 2014). Precise identification of IDC is required to direct patient treatment. In recent years, many machine learning models using deep neural networks have been used to aid medical and biological classification (Janowczyk and Madabhushi 2016; Cruz-Roa et al. 2014). Thus, one potential benefit of a predictive model for medical image classification is to speed-up the diagnosis process and begin treatment for patients.

A personal motivation to pursue this project was to get exposure working with image data. I anticipated the following challenges. First, reading and storing image data. Unlike rectangular datasets, each image was a .jpg file and the entire data set was spread across nested directories. Moreover, each image is a big file and loading all of them together was computationally intensive. Second, the quality of an image is contingent on many factors, eg., the device from which it is taken, the magnification, the angle, blur, zoom, flip, etc. Therefore, two images that appear the same to human eye, might be very encoded differently by the computer. To overcome this challenge, data augmentation is performed to add random noise to the training dataset for replicating close to real-life scenarios. However, doing so would increase the sample size, which further burdens the computational resources. In this project, however, since I am working with medical images taken by professional radiologists, I do not expect image quality to be a major roadblock. Lastly, I am working with colored images, therefore, each image will be represented by a three-dimensional array, with each cell in the matrix representing a pixel. Though, in my professional life, I don't anticipate working with medical or biological images, I believe this is a good first step to building skills for image analysis using machine learning.

Description of the data

The section on description of the data is divided into the following subsections: data source, data representation, feature extraction, and data reduction.

Data Source

The data comes from a 2014 study (Cruz-Roa et al. 2014), with images of the entire region of the breast for 162 cancer patients. I accessed it through Kaggle (link provided in the beginning of the report). It consisted of whole mount slide images of Breast Cancer specimens scanned at 40x. Each slide, covering the entire breast area, was broken into 277,524 patches of size 50 x 50 in RGB. Of the total, 198,738 (71.5%) patches were IDC negative and 78,786 (28.5%) patches were IDC positive.

Data Representation

The data were provided in .jpg format. I used Python module Keras (Chollet 2015) to preprocess the images. First, I read the .jpg image in the environment. Next, I converted each image to an array. As a result, each patch of image was represented by a three-dimensional array of 50 x 50, where each value in the matrix represents the pixel brightness on the particular row and column combination. The dimensions represents the red, blue, and green light channels (RGB). Figure 1 shows how one image looked in different light channels.

(a) Image in RGB channel   (c) Image in Red channel

(b) Image in Blue Channel   (d) Image in Green channel

Figure 1: Image in different light channels

Figure 2 shows five random images in RGB and whether the patch was cancerous or not (IDC positive or negative).
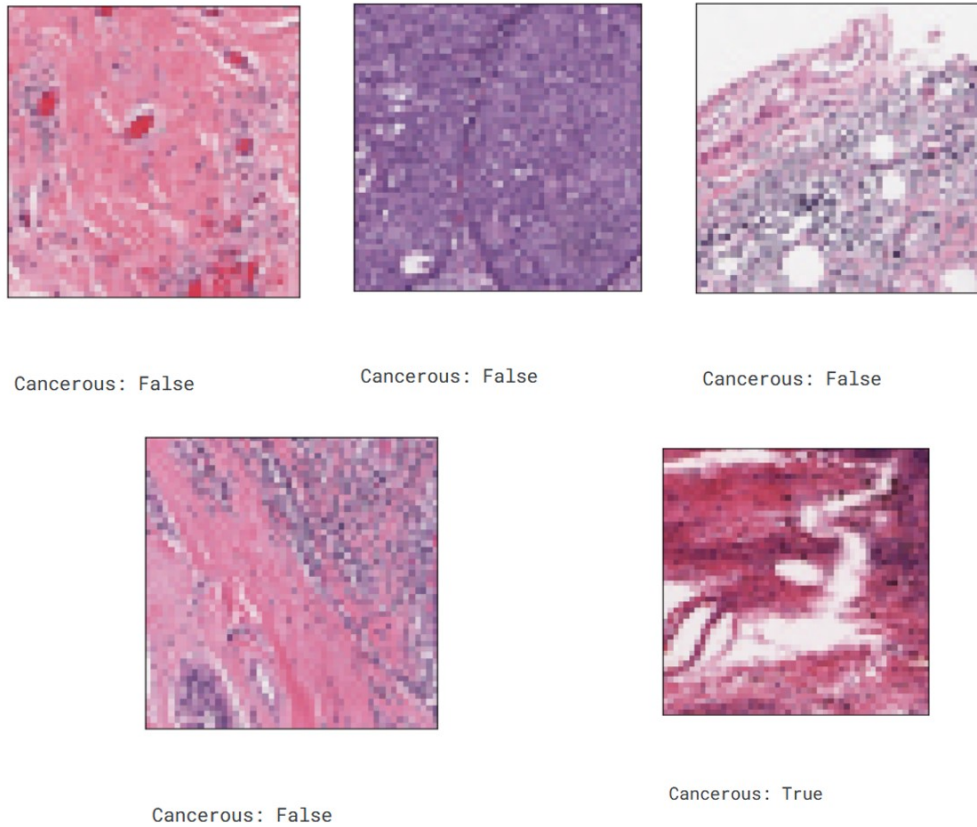

Cancerous: False   Cancerous: False   Cancerous: False

Cancerous: False

Cancerous: True

Figure 2: Random Sample of 5 images

Thus, each image file can be associated with 7,500 (50 x 50 x 3) numbers representing the pixel brightness in the three light channels. However, these number are not meaningful in themselves, as they did not communicate how the image appears as a whole. There are several ways image features can be calculated, for instance, the mean intensity of image in all channels, the edges and boundaries in an image, number of elements in an image, etc. Since it required expertise to know which features to extract, I decided not to hand-define any custom features. Instead, I used a pre-trained image classification model, ResNet-50, which generated 8,192 features for the images in this data set.

Data Features using ResNet-50

ResNet-50, short for Residual Networks, is a convolutional neural network that is 50 layers deep. It has been pre-trained on more than a million images from the ImageNet database, and one if its more complex architecture won the ImageNet competition in 2015 (He et al. (2016)). The pre-trained network can classify images into 1000 object categories, such as keyboard, mouse, pencil, and many animals. The network architecture is shown in Figure 3.

I used Python module, Keras(Chollet 2015) and TensorFlow(Abadi et al. 2015), to load the pre-trained model with original training weights and extracted 8,142 features from each image.
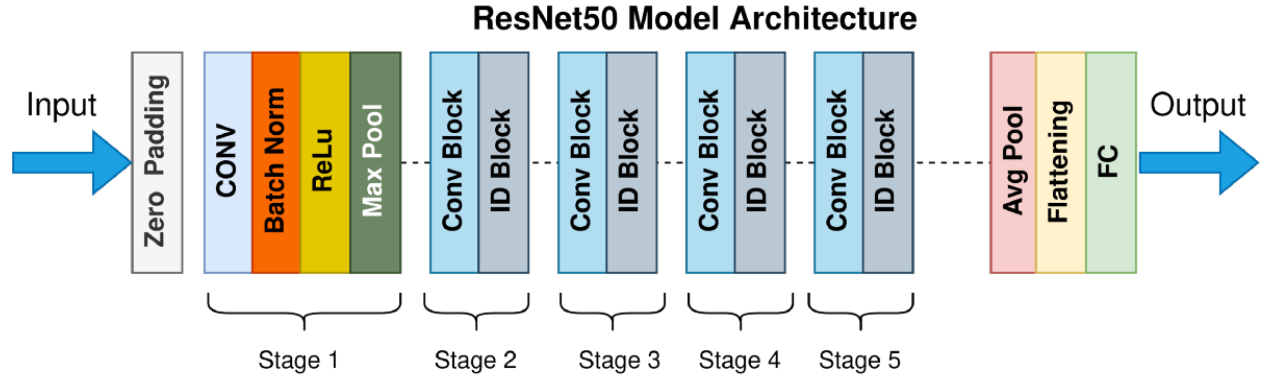
Figure 3: ResNet-50 architecture

Data Reduction

Given the very large size of the original dataset and the limited computations power on my end, I decided to work with two subsets samples of this data, one relatively large sample with close to 50,000 random images, and other, a smaller sample with 5,000 images. The smaller sample was taken from the larger sample.

The larger sample was a random selection of close to 50,000 images. The dataset with all 8,000 features was 3Gb large and took up 20Gb RAM on loading. Since the size drastically limited the resources for model training, I decided to reduce the features for this sample. First, I splitted the data into training and testing datasets. Next, on the whole training dataset, I removed all features that had zero variance, near zero variance, or a correlation of 0.5 or higher. This reduced the number of features to 1,371. Thus, the larger sample had a final size of 49,921 rows and 1,373 columns.

The smaller sample had all the 8,192 features. However, the recipe included steps to remove columns based on the same criteria as mentioned above. This dataset had a final size of 5,004 rows and 8,194 columns. Note there was one ID column (serian_num) and one outcome column (truth_val) in both the datasets.

Description of the models

At heart, this was a classification problem with two classes, positive or negative for IDC. Therefore, I decided to use logistic regression. Given the large number of features, over-fitting was possible, and so I also fit penalized logistic regressions models with both ridge penalty and lasso penalty. Logistic regression is analogous to linear regression. However, instead of predicting a continuous outcome, it predicts the probability of an observation to be in one of the two outcome classes. Model selection was done using the logistic loss criteria, i.e., a model with the small logistic loss was selected.

Logistic regression assumes that the predictors have a linear relationship with the outcome, all important predictors are present in the model, and there is no multicollinearity. In the regularized version of the logistic regression, we penalize the model for over-fitting by adding a penalty term. This constrains the size of the predicted coefficients. In unregularized logistic regression, there are no hyperparameter to be tuned. In the ridge regression, the penalty is the euclidean norm, whereas in lasso regression, the penalty is the least absolute shrinkage. Thus, there was one hyperparameter for each ridge and lasso regressions to tune.

Implementation Plan

First, split the sample into training and test set, with an 80-20 ratio. Next, make indices for 10-fold cross validation. Then, use recipe package, and define the first column (serial_num) as an ID, the second column

(truth_val) as an outcome, and the remaining columns as predictors. Next, remove any variable with zero variance or near-zero variance, and keep only one of the columns which are medium to highly correlated (a correlation of 0.5 or more). Lastly, transform the outcome into 1 and 2 instead of 0 and 1, to avoid a run-time error.

Metric for model comparison

For model comparison, I will use logistic loss (logLoss) and area under the receiver operating characteristic (ROC) curve (AUC) as the metric. I will also compute the confusion metric, and evaluate the true positive rate, true negative rate, accuracy, and precision.

Packages required

I used the following Python and R packages for the project: pandas (team 2020), NumPy (Harris et al. 2020), keras (Chollet 2015), scikit-leanr (Pedregosa et al. 2011), TensorFlow (Abadi et al. 2015), Matplotlib (Hunter 2007), tidyverse (Wickham et al. 2019), recipes (Kuhn and Wickham 2022), caret (Kuhn 2022), base R (R Core Team 2021), cutpointr(Thiele and Hirschfeld 2021), vip(Greenwell and Boehmke 2020), glmnet(Friedman, Hastie, and Tibshirani 2010).

Model Fit

I will discuss the model results for the small and large samples. While discussion the results, I will share the metrics for model comparison, the classification plot, as well as the logistic loss plot for the hyperparamters.

Model Fit for small sample

Table 1

Model performance for small sample

Model	logLoss	Cut-off	AUC	Accuracy	True Positive rate	True Negative rate	Precision
Unregularized Logistic Regression	0.371	0.286	0.914	0.858	0.849	0.86	0.686
Logistic Regression with Ridge Penalty	0.342	0.297	0.921	0.867	0.857	0.87	0.703
Logistic Regression with Lasso Penalty	0.357	0.308	0.919	0.861	0.834	0.87	0.697

Table 1 summarizes the model fit on the small sample for the three models. The logLoss for unregularized, regularized with ridge penalty, and regularized with lasso penalty were 0.371, 0.342, and 0.357 respectively. The area under the ROC curve for the three models were 0.914, 0.921, and 0.919. As was expected,

regularization improved the results slightly, as this sample had relatively large features to observation ratio. The best tune hyperparameters for the ridge logistic regression was 0.066, and for lasso losgistic regression was 0.0038. On this smaller sample, I choose the regularized logistic regression with ridge penalty as the final model for it's slight improvement. I chose a cut-off of 0.297 for this model. This cut-off was found found using the cutpointr package (Thiele and Hirschfeld 2021), and maximizes the distinction between the negative and positive sets. Figures 4-6 shows the classification plots for the three models on the small sample.
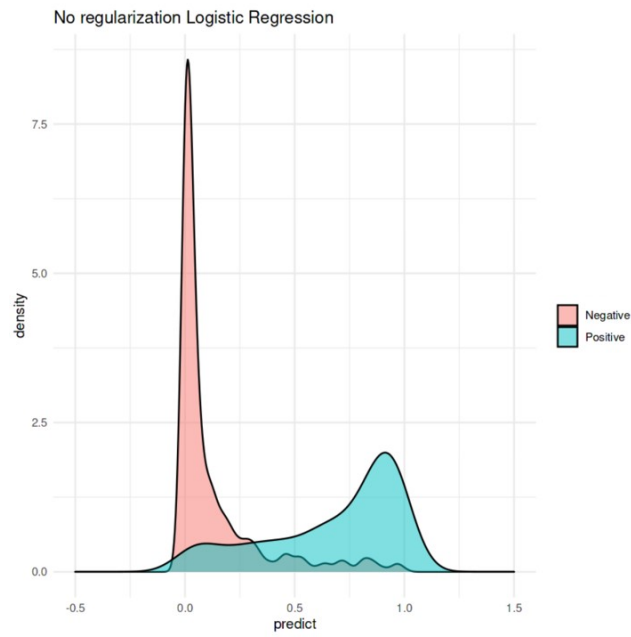

Figure 4: Classification Plot for unregularized logistic regression plot

Figure 5: Regularization paramter and classification plot for regularized logistic regression (ridge penalty)
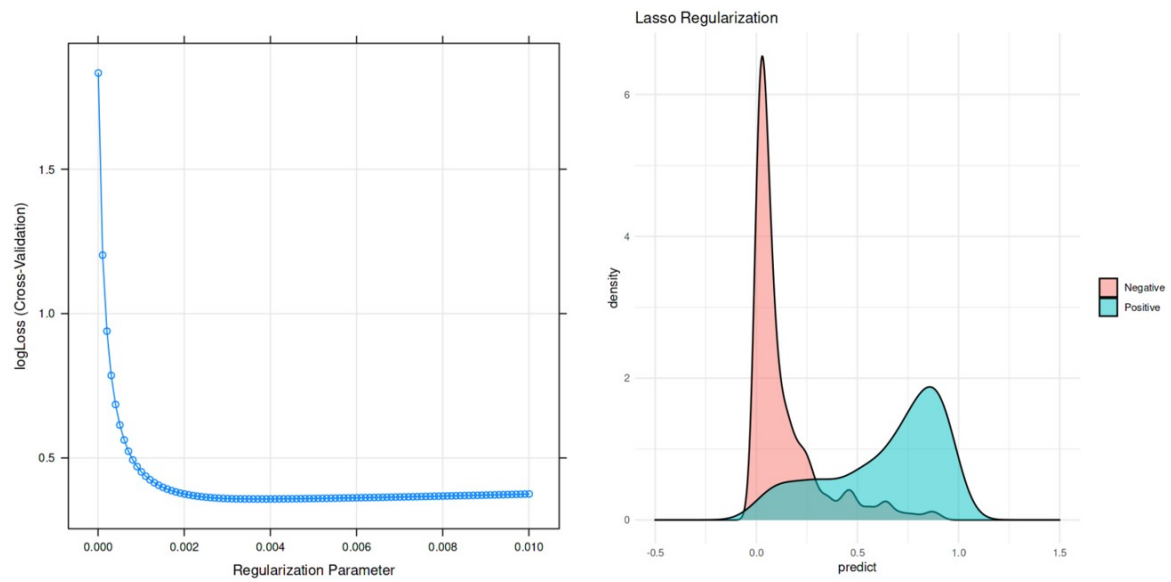

Figure 6: Regularization paramter and classification plot for regularized logistic regression (lasso penalty)

Model Fit for large sample

Table 2 summarizes the model fit on the small sample for the three models. The logLoss for unregularized, regularized with ridge penalty, and regularized with lasso penalty were 0.369, 0.368, and 0.369 respectively. The area under the ROC curve for the three models was 0.897 in all cases. As was expected, all regressions performed alike becuase of the smaller features to observations ratio. The best tune hyperparameters for the ridge logistic regression was 0.031, and for lasso losgistic regression was 0.001. On this larger sample, I choose the unregularized logistic regression for slightly higher precision. I chose a cut-off of 0.283 for this model. This cut-off was found found using the cutpointr package (Thiele and Hirschfeld 2021), and maximizes the distinction between the negative and positive sets. Figures 7-9 show classification plots for the three models for the larger dataset.

Table 2

Model performance for large sample

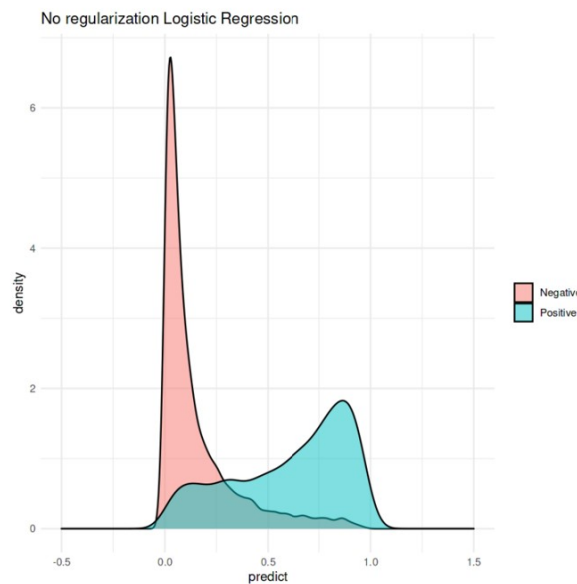| Model	logLoss	Cut-off	AUC	Accuracy	True Positive rate	True Negative rate	Precision
Unregularized Logistic Regression	0.369	0.283	0.897	0.83	0.824	0.832	0.657
Logistic Regression with Ridge Penalty	0.368	0.275	0.897	0.827	0.832	0.825	0.65
Logistic Regression with Lasso Penalty	0.36	0.274	0.897	0.826	0.832	0.823	0.647

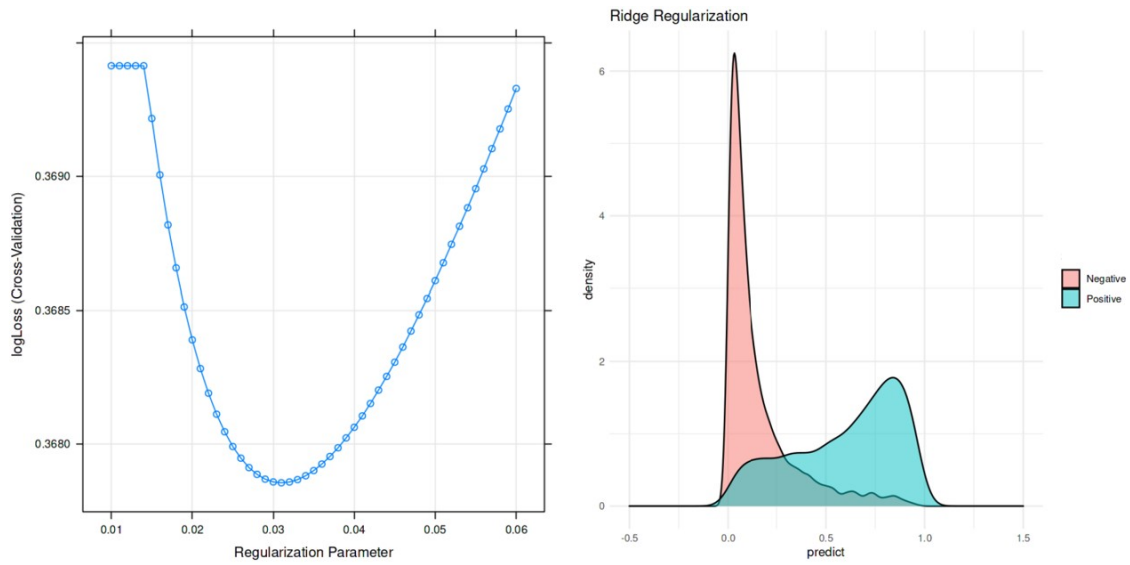Figure 7: Classification Plot for unregularized logistic regression plot

Figure 8: Regularization paramter and classification plot for regularized logistic regression (ridge penalty)
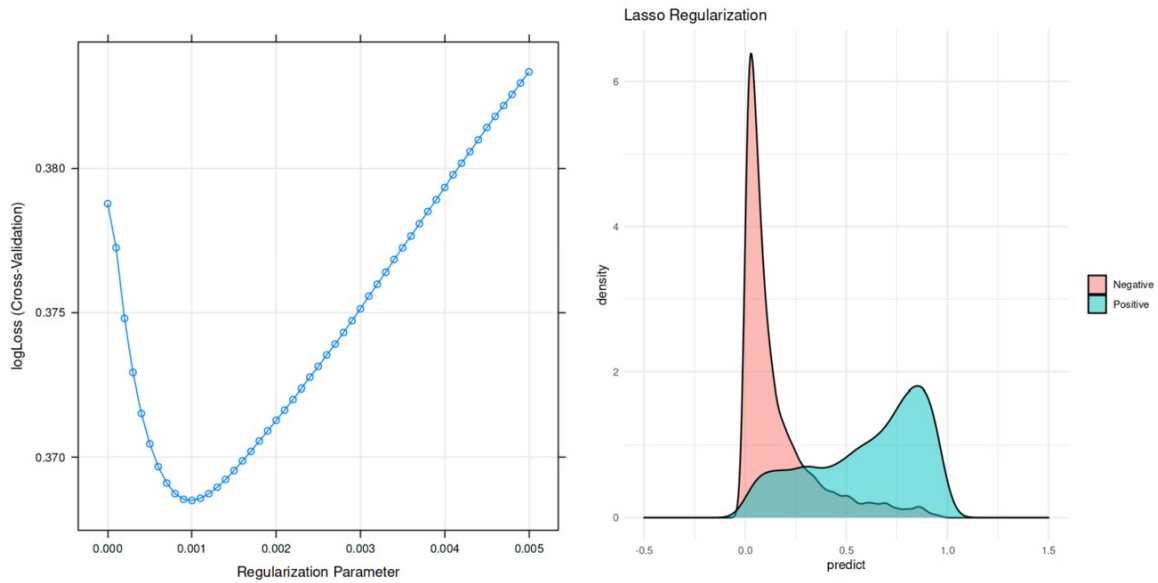

Figure 9: Regularization paramter and classification plot for regularized logistic regression (lasso penalty)

Discussion

All the three logistic models performed similarly, and the AUC was around 0.9 in all models and samples sizes. This indicates that regularization doesn't significantly improve results. One reason for this could be that I removed moderate to highly correlated columns in the beginning, bringing the features to less than 1/4th of its original size. This could have reduced the possibility of potential over-fit. It would be interesting to see how the models performed without removing correlated columns, or removing only those columns with high correlation (0.8 or higher).

In the smaller sample, however, there were small improvements with regularization (a 0.1 improvement in the AUC metric). This could be because there were about 1,000 features for a training set of 4,000, i.e., a features to observation ratio of 1:4. In the larger sample, there was no difference between the models. Here, there were 1,300 features but 40,000 observation, a feature to observation ratio of close to 1:40. Overall, and to my surprise, the models did well and had an accuracy of around 83-86%. It would be interesting to compare the results with the original 7,500 pixel features and compare it to the current model, to know whether the model would still perform decently without the aid of ResNet-50 features. Another point to note was that the model performance slightly dropped on the larger set, this could probably be because of the larger sample size.

In terms of importance, there were around 20 ResNet-50 features that were assigned an importance of more than 50, and 700 features with importance of more than 20. However, just reading the features wasn't very insightful by itself. There are a number of limitations of this project. First, the feature-generating model, ResNet-50, was not developed on medical images. Second, due to the high computation requirement for the projects, I could only work with a subset of the total sample. Next, I did not perform any data augmentation on the images, something that is plausible in real-life scenarios. Lastly, though the simple logistic regression model is comparatively easier to explain and led to acceptable performance when built on top of ResNet-50 features, classical machine learning approaches are rarely used in modern image analysis problems for a number of reasons. Thus, future models could be built using neural networks approaches or model ensembles.

On my personal end, there were a number of learning points during the project. Reading and pre-processing image data was something very new I tried. To be able to run models, I had to make a number of decision - how much sample to consider, when to reduce size of the data, etc., and this made me cognizant of the decisions one has to make in machine learning practices.

References

Abadi, Martín, Ashish Agarwal, Paul Barham, Eugene Brevdo, Zhifeng Chen, Craig Citro, Greg S. Corrado, et al. 2015. "TensorFlow: Large-Scale Machine Learning on Heterogeneous Systems." https://www.tensorflow.org/.

Chollet, François. 2015. "Keras." *GitHub Repository.* https://github.com/fchollet/keras; GitHub.

Cruz-Roa, Angel, Ajay Basavanhally, Fabio González, Hannah Gilmore, Michael Feldman, Shridar Ganesan, Natalie Shih, John Tomaszewski, and Anant Madabhushi. 2014. "Automatic Detection of Invasive Ductal Carcinoma in Whole Slide Images with Convolutional Neural Networks." In *Medical Imaging 2014: Digital Pathology*, 9041:904103. SPIE.

Ferlay, Jacques, Murielle Colombet, Isabelle Soerjomataram, Donald M Parkin, Marion Piñeros, Ariana Znaor, and Freddie Bray. 2021. "Cancer Statistics for the Year 2020: An Overview." *International Journal of Cancer* 149 (4): 778–89.

Friedman, Jerome, Trevor Hastie, and Robert Tibshirani. 2010. "Regularization Paths for Generalized Linear Models via Coordinate Descent." *Journal of Statistical Software* 33 (1): 1–22. https://doi.org/10.18637/jss.v033.i01.

Greenwell, Brandon M., and Bradley C. Boehmke. 2020. "Variable Importance Plots—an Introduction to the Vip Package." *The R Journal* 12 (1): 343–66. https://doi.org/10.32614/RJ-2020-013.

Harris, Charles R., K. Jarrod Millman, Stéfan J. van der Walt, Ralf Gommers, Pauli Virtanen, David Cournapeau, Eric Wieser, et al. 2020. "Array Programming with NumPy." *Nature* 585 (7825): 357–62. https://doi.org/10.1038/s41586-020-2649-2.

He, Kaiming, Xiangyu Zhang, Shaoqing Ren, and Jian Sun. 2016. "Deep Residual Learning for Image Recognition." In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, 770–78.

Heron, Melonie. 2021. "National Vital Statistics Reports." *National Vital Statistics Reports* 70 (3).

Hunter, J. D. 2007. "Matplotlib: A 2d Graphics Environment." *Computing in Science & Engineering* 9 (3): 90–95. https://doi.org/10.1109/MCSE.2007.55.

Janowczyk, Andrew, and Anant Madabhushi. 2016. "Deep Learning for Digital Pathology Image Analysis: A Comprehensive Tutorial with Selected Use Cases." *Journal of Pathology Informatics* 7 (1): 29.

Kuhn, Max. 2022. *Caret: Classification and Regression Training.* https://CRAN.R-project.org/package=caret.

Kuhn, Max, and Hadley Wickham. 2022. *Recipes: Preprocessing and Feature Engineering Steps for Modeling.* https://CRAN.R-project.org/package=recipes.

Pedregosa, F., G. Varoquaux, A. Gramfort, V. Michel, B. Thirion, O. Grisel, M. Blondel, et al. 2011. "Scikit-Learn: Machine Learning in Python." *Journal of Machine Learning Research* 12: 2825–30.

R Core Team. 2021. *R: A Language and Environment for Statistical Computing.* Vienna, Austria: R Foundation for Statistical Computing. https://www.R-project.org/.

Sharma, Ganesh N, Rahul Dave, Jyotsana Sanadya, Piush Sharma, KK3255438 Sharma, et al. 2010. "Various Types and Management of Breast Cancer: An Overview." *Journal of Advanced Pharmaceutical Technology & Research* 1 (2): 109.

team, The pandas development. 2020. *Pandas-Dev/Pandas: Pandas* (version latest). Zenodo. https://doi.org/10.5281/zenodo.3509134.

Thiele, Christian, and Gerrit Hirschfeld. 2021. "cutpointr: Improved Estimation and Validation of Optimal Cutpoints in R." *Journal of Statistical Software* 98 (11): 1–27. https://doi.org/10.18637/jss.v098.i11.

Wickham, Hadley, Mara Averick, Jennifer Bryan, Winston Chang, Lucy D'Agostino McGowan, Romain François, Garrett Grolemund, et al. 2019. "Welcome to the tidyverse." *Journal of Open Source Software* 4 (43): 1686. https://doi.org/10.21105/joss.01686.